

Task 2. Annual Salary Prediction

July 31, 2020

1 Task 2 - Predictive Analytics

For the Data@ANZ Virtual Experience Program

1. Build a simple regression model to predict the annual salary for each customer using the attributes you identified above

How accurate is your model? Should ANZ use it to segment customers (for whom it does not have this data) into income brackets for reporting purposes?

2. For a challenge: build a decision-tree based model to predict salary. Does it perform better? How would you accurately test the performance of this model?

```
[1]: import pandas as pd
import numpy as np
```

```
[2]: df = pd.read_excel('ANZ synthesised transaction dataset.xlsx')
```

```
[3]: #df.info()
```

1.1 Explore salary feature

1.1.1 First we identify monthly salary for each customer

```
[4]: df['month'] = pd.DatetimeIndex(df['date']).month
df[['customer_id', 'date', 'month']]
```

```
[4]:
```

	customer_id	date	month
0	CUS-2487424745	2018-08-01	8
1	CUS-2487424745	2018-08-01	8
2	CUS-2142601169	2018-08-01	8
3	CUS-1614226872	2018-08-01	8
4	CUS-2487424745	2018-08-01	8
...
12038	CUS-55310383	2018-10-31	10
12039	CUS-2688605418	2018-10-31	10
12040	CUS-2663907001	2018-10-31	10
12041	CUS-1388323263	2018-10-31	10
12042	CUS-3129499595	2018-10-31	10

[12043 rows x 3 columns]

```
[5]: df['txn_description'].value_counts()
```

```
[5]: SALES-POS      3934
      POS          3783
      PAYMENT      2600
      PAY/SALARY    883
      INTER BANK   742
      PHONE BANK   101
      Name: txn_description, dtype: int64
```

```
[6]: df_salary_payment = df[df['txn_description']=='PAY/SALARY'].copy()
      df_salary_payment.sort_values(by=['customer_id'], inplace=True)
      df_salary_payment.head(5)
```

```
[6]:      status  card_present_flag  bpay_biller_code      account  currency \
2530  posted                    NaN                0  ACC-2828321672      AUD
4402  posted                    NaN                0  ACC-2828321672      AUD
8142  posted                    NaN                0  ACC-2828321672      AUD
1744  posted                    NaN                0  ACC-2828321672      AUD
6271  posted                    NaN                0  ACC-2828321672      AUD
```

```
      long_lat  txn_description  merchant_id  merchant_code  first_name \
2530  153.03 -27.51      PAY/SALARY        NaN            0.0  Stephanie
4402  153.03 -27.51      PAY/SALARY        NaN            0.0  Stephanie
8142  153.03 -27.51      PAY/SALARY        NaN            0.0  Stephanie
1744  153.03 -27.51      PAY/SALARY        NaN            0.0  Stephanie
6271  153.03 -27.51      PAY/SALARY        NaN            0.0  Stephanie
```

```
      ...  merchant_suburb  merchant_state      extraction \
2530  ...                NaN            NaN  2018-08-21T16:00:00.000+0000
4402  ...                NaN            NaN  2018-09-04T16:00:00.000+0000
8142  ...                NaN            NaN  2018-10-02T16:00:00.000+0000
1744  ...                NaN            NaN  2018-08-14T16:00:00.000+0000
6271  ...                NaN            NaN  2018-09-18T16:00:00.000+0000
```

```
      amount      transaction_id      country      customer_id \
2530  970.47  71cd874fc20741f8b4a589c8286afeb2  Australia  CUS-1005756958
4402  970.47  e588bd113b3645ee82fb386e336c42a1  Australia  CUS-1005756958
8142  970.47  6a0796f6e44c4d49b288a593bdc23503  Australia  CUS-1005756958
1744  970.47  deaff82de78840f08a035e5404ce5e29  Australia  CUS-1005756958
6271  970.47  6b622e0b12324ac2a1b6c946f43bce04  Australia  CUS-1005756958
```

```
      merchant_long_lat  movement  month
2530                NaN      credit      8
```

4402	NaN	credit	9
8142	NaN	credit	10
1744	NaN	credit	8
6271	NaN	credit	9

[5 rows x 24 columns]

```
[7]: df_salary_payment_details = df_salary_payment[['customer_id', 'month', 'amount']]
df_customer_salary = df_salary_payment_details.copy()
```

```
[8]: # Calculate customer salary payment per month
df_customer_salary = df_salary_payment_details.groupby(by=['customer_id', 'month'], as_index=False).sum()
df_customer_salary.rename(columns={'amount': 'amount_month'}, inplace=True)
df_customer_salary.head(6)
```

```
[8]:
```

	customer_id	month	amount_month
0	CUS-1005756958	8	3881.88
1	CUS-1005756958	9	3881.88
2	CUS-1005756958	10	4852.35
3	CUS-1117979751	8	7157.30
4	CUS-1117979751	9	7157.30
5	CUS-1117979751	10	10735.95

```
[9]: # Calculate customer number of salary payments per month
df_num_salary = df_salary_payment_details.groupby(by=['customer_id', 'month'], as_index=False).count()
df_customer_salary['num_of_payments'] = df_num_salary['amount']
df_customer_salary.head(6)
```

```
[9]:
```

	customer_id	month	amount_month	num_of_payments
0	CUS-1005756958	8	3881.88	4
1	CUS-1005756958	9	3881.88	4
2	CUS-1005756958	10	4852.35	5
3	CUS-1117979751	8	7157.30	2
4	CUS-1117979751	9	7157.30	2
5	CUS-1117979751	10	10735.95	3

```
[10]: # Calculate customer base salary payment
df_base_salary = df_salary_payment_details.groupby(by=['customer_id', 'month'], as_index=False).agg(np.average)
df_customer_salary['base_salary'] = df_base_salary['amount']
df_customer_salary.head(6)
```

```
[10]:
```

	customer_id	month	amount_month	num_of_payments	base_salary
0	CUS-1005756958	8	3881.88	4	970.47
1	CUS-1005756958	9	3881.88	4	970.47

2	CUS-1005756958	10	4852.35	5	970.47
3	CUS-1117979751	8	7157.30	2	3578.65
4	CUS-1117979751	9	7157.30	2	3578.65
5	CUS-1117979751	10	10735.95	3	3578.65

1.1.2 Calculate Annual Salary Payment per Customer

```
[11]: df_annual_salary = df_salary_payment_details.drop(columns=['month'])
df_annual_salary = df_annual_salary.groupby(['customer_id'], as_index = False).
    →sum()
df_annual_salary.rename(columns={'amount': 'annual_salary'}, inplace=True)
df_annual_salary.head(6)
```

```
[11]:      customer_id  annual_salary
0  CUS-1005756958      12616.11
1  CUS-1117979751      25050.55
2  CUS-1140341822      11499.06
3  CUS-1147642491      22248.07
4  CUS-1196156254      27326.11
5  CUS-1220154422      15976.52
```

```
[12]: df_total_payments = df_customer_salary[['customer_id', 'num_of_payments']].
    →groupby(['customer_id'], as_index=False).sum()
df_annual_salary['num_payments'] = df_total_payments['num_of_payments']
df_annual_salary.head(6)
```


```
[12]:      customer_id  annual_salary  num_payments
0  CUS-1005756958      12616.11           13
1  CUS-1117979751      25050.55            7
2  CUS-1140341822      11499.06            6
3  CUS-1147642491      22248.07           13
4  CUS-1196156254      27326.11            7
5  CUS-1220154422      15976.52            7
```

```
[13]: df_annual_salary['avg_num_payments'] = df_annual_salary['num_payments'].
    →apply(lambda x : round(x/3.))
df_annual_salary.head(6)
```

```
[13]:      customer_id  annual_salary  num_payments  avg_num_payments
0  CUS-1005756958      12616.11           13             4
1  CUS-1117979751      25050.55            7             2
2  CUS-1140341822      11499.06            6             2
3  CUS-1147642491      22248.07           13             4
4  CUS-1196156254      27326.11            7             2
5  CUS-1220154422      15976.52            7             2
```

```
[14]: df_annual_salary['annual_salary'].hist(bins=10)
```


```
[14]: <matplotlib.axes._subplots.AxesSubplot at 0x1d3f4100470>
```



Approximately 3 people are the ones with the highest payment salaries.

```
[15]: df_annual_salary['avg_num_payments'].hist(bins=10)
```

```
[15]: <matplotlib.axes._subplots.AxesSubplot at 0x1d3f486e470>
```



A little more than half of 100 customers have 2 payments per month. And one third has 4 payments per month. Could be an important feature to take into consideration.

```
[16]: df_annual_salary[df_annual_salary['customer_id'] == 'CUS-1005756958']
```

```
[16]:      customer_id  annual_salary  num_payments  avg_num_payments
0  CUS-1005756958      12616.11           13              4
```

```
[17]: df_customer_salary[df_customer_salary['customer_id'] == 'CUS-1005756958']
```

```
[17]:      customer_id  month  amount_month  num_of_payments  base_salary
0  CUS-1005756958      8      3881.88              4      970.47
1  CUS-1005756958      9      3881.88              4      970.47
2  CUS-1005756958     10      4852.35              5      970.47
```

- The relevant fields in “df_annual_salary” are ‘amount’ and ‘avg_num_payments’
- The relevant field in “df_customer_salary” is ‘salary’

1.2 Explore correlations between annual salary and various customer attributes

There are some original features(readily available in the data) that can be relevant to compare against annual salary, such as: ‘age’, ‘gender’, ‘long_lat’ and ‘balance’

```
[18]: df_salary_payment[['customer_id', 'age', 'gender', 'long_lat', 'balance']].nunique()
```

```
[18]: customer_id    100
      age           33
      gender        2
      long_lat      100
      balance       883
      dtype: int64
```

It seems the feature BALANCE has different values in each customer. It needs to be average per each customer.

```
[19]: # Calculating the average of balance for each customer.
      df_avg_balance = df_salary_payment[['customer_id', 'balance']].
      →groupby(['customer_id'], as_index=False).mean()
      df_avg_balance.head(6)
```

```
[19]:      customer_id      balance
0  CUS-1005756958  4718.665385
1  CUS-1117979751  11957.202857
2  CUS-1140341822   5841.720000
3  CUS-1147642491   8813.467692
4  CUS-1196156254  23845.717143
5  CUS-1220154422   9225.907143
```

1.3 Get unique data per customer

```
[20]: df_final = df_salary_payment[['customer_id', 'age', 'gender', 'long_lat']].copy()
      df_final.drop_duplicates('customer_id', inplace=True)
      df_final.reset_index(drop=True, inplace=True)
      df_final.head(6)
```

```
[20]:      customer_id  age  gender      long_lat
0  CUS-1005756958   53      F  153.03 -27.51
1  CUS-1117979751   21      M  115.81 -31.82
2  CUS-1140341822   28      M  144.97 -37.42
3  CUS-1147642491   34      F  151.04 -33.77
4  CUS-1196156254   34      F  138.52 -35.01
5  CUS-1220154422   25      F  150.50 -23.40
```

```
[21]: df_final['balance'] = df_avg_balance['balance']
```

Features like GENDER and LOCATION(long_lat) need to be converted to numbers in order to analyze the correlation between salary.

```
[22]: # convert gender column to Integer type
      df_final['gender'] = df['gender'].map( {'M':1, 'F':0} )
      df_final.head(6)
```

```
[22]:
```

	customer_id	age	gender	long_lat	balance
0	CUS-1005756958	53	0	153.03 -27.51	4718.665385
1	CUS-1117979751	21	0	115.81 -31.82	11957.202857
2	CUS-1140341822	28	1	144.97 -37.42	5841.720000
3	CUS-1147642491	34	0	151.04 -33.77	8813.467692
4	CUS-1196156254	34	0	138.52 -35.01	23845.717143
5	CUS-1220154422	25	1	150.50 -23.40	9225.907143

Calculate distance of customers from Centre of Australia

If we analyze long and lat as single values, that probably won't tell us too much things.

So we are going to use those values to calculate how far a customer is from the centre of Australia (Lambert Gravitational Centre) [latitude: -25.610111, longitude: 134.354806]

Reference: https://www.wikiwand.com/en/Centre_points_of_Australia

```
[23]: long_lat = df_final['long_lat'].str.split(" ", n = 1, expand = True)
df_final['long'] = long_lat[0].astype('float')
df_final['lat'] = long_lat[1].astype('float')
df_final.drop(['long_lat'], axis = 1, inplace=True)
df_final.head(6)
```

```
[23]:
```

	customer_id	age	gender	balance	long	lat
0	CUS-1005756958	53	0	4718.665385	153.03	-27.51
1	CUS-1117979751	21	0	11957.202857	115.81	-31.82
2	CUS-1140341822	28	1	5841.720000	144.97	-37.42
3	CUS-1147642491	34	0	8813.467692	151.04	-33.77
4	CUS-1196156254	34	0	23845.717143	138.52	-35.01
5	CUS-1220154422	25	1	9225.907143	150.50	-23.40

```
[24]: centre_latitude = -25.610111
centre_longitude = 134.354806

import math
def calculateDistance(row):
    return math.sqrt((row['long'] - centre_longitude)**2 + (row['lat'] -
    ↪centre_latitude)**2)
```

```
[25]: df_final['location_dist'] = df_final.apply(lambda row: calculateDistance(row),
    ↪axis=1)
df_final.drop(['long', 'lat'], axis = 1, inplace=True)
df_final.head(6)
```

```
[25]:
```

	customer_id	age	gender	balance	location_dist
0	CUS-1005756958	53	0	4718.665385	18.771586
1	CUS-1117979751	21	0	11957.202857	19.556905
2	CUS-1140341822	28	1	5841.720000	15.879415
3	CUS-1147642491	34	0	8813.467692	18.573623

4	CUS-1196156254	34	0	23845.717143	10.281379
5	CUS-1220154422	25	1	9225.907143	16.295763

Now, we can concatenate the final information with values we derived/construct ourselves.

- from “df_annual_salary” -> ‘amount’(total salary over months) and ‘avg_num_payments’
- from “df_customer_salary” -> ‘salary’ (base salary)

```
[26]: df_customer_unique = df_customer_salary.drop_duplicates(['customer_id']).
      ↪reset_index(drop = True)
```

1.3.1 Concatenate original selected features with created ones

```
[27]: data_annual = pd.concat([df_final, df_annual_salary[['annual_salary',
      ↪'avg_num_payments']]], axis=1)
data = pd.concat([data_annual, df_customer_unique['base_salary']], axis =1 )
data.head(6)
```

```
[27]:
```

	customer_id	age	gender	balance	location_dist	annual_salary \
0	CUS-1005756958	53	0	4718.665385	18.771586	12616.11
1	CUS-1117979751	21	0	11957.202857	19.556905	25050.55
2	CUS-1140341822	28	1	5841.720000	15.879415	11499.06
3	CUS-1147642491	34	0	8813.467692	18.573623	22248.07
4	CUS-1196156254	34	0	23845.717143	10.281379	27326.11
5	CUS-1220154422	25	1	9225.907143	16.295763	15976.52

	avg_num_payments	base_salary
0	4	970.47
1	2	3578.65
2	2	1916.51
3	4	1711.39
4	2	3903.73
5	2	2282.36

```
[28]: # reorder columns
columns = [col for col in data.columns if col != 'annual_salary'] +
      ↪['annual_salary']
data = data[columns]
data.head(6)
```

```
[28]:
```

	customer_id	age	gender	balance	location_dist	avg_num_payments \
0	CUS-1005756958	53	0	4718.665385	18.771586	4
1	CUS-1117979751	21	0	11957.202857	19.556905	2
2	CUS-1140341822	28	1	5841.720000	15.879415	2
3	CUS-1147642491	34	0	8813.467692	18.573623	4
4	CUS-1196156254	34	0	23845.717143	10.281379	2
5	CUS-1220154422	25	1	9225.907143	16.295763	2

	base_salary	annual_salary
0	970.47	12616.11
1	3578.65	25050.55
2	1916.51	11499.06
3	1711.39	22248.07
4	3903.73	27326.11
5	2282.36	15976.52

1.3.2 Calculate Pearson Correlation between fields

```
[29]: df_all_corr = data.corr()
df_all_corr
```

```
[29]:
```

	age	gender	balance	location_dist	\
age	1.000000	-0.021263	0.227026	-0.087073	
gender	-0.021263	1.000000	0.059663	-0.100321	
balance	0.227026	0.059663	1.000000	-0.033805	
location_dist	-0.087073	-0.100321	-0.033805	1.000000	
avg_num_payments	0.187163	-0.051730	-0.192136	-0.061067	
base_salary	-0.135264	-0.072938	0.231019	0.070881	
annual_salary	-0.036504	-0.109313	0.198755	0.097938	

	avg_num_payments	base_salary	annual_salary
age	0.187163	-0.135264	-0.036504
gender	-0.051730	-0.072938	-0.109313
balance	-0.192136	0.231019	0.198755
location_dist	-0.061067	0.070881	0.097938
avg_num_payments	1.000000	-0.693218	-0.030318
base_salary	-0.693218	1.000000	0.534883
annual_salary	-0.030318	0.534883	1.000000

```
[30]: df_all_corr= df_all_corr.abs().unstack().sort_values(kind="quicksort",
→ascending=False).reset_index()
df_all_corr.rename(columns={"level_0": "Feature 1", "level_1": "Feature 2", 0:
→'Correlation Coefficient'}, inplace=True)
df_all_corr[df_all_corr['Feature 1'] == 'annual_salary']
```

```
[30]:
```

	Feature 1	Feature 2	Correlation Coefficient
0	annual_salary	annual_salary	1.000000
10	annual_salary	base_salary	0.534883
15	annual_salary	balance	0.198755
23	annual_salary	gender	0.109313
28	annual_salary	location_dist	0.097938
41	annual_salary	age	0.036504
46	annual_salary	avg_num_payments	0.030318

“base_salary” has a strong correlation with annual_salary.

We can not conclude changes in the variable cause changes in annual_salary based on correlation alone. Only properly controlled experiments enable us to determine whether a relationship is causal.

A low Pearson correlation coefficient does not mean that no relationship exists between the variables.

The variables may have a nonlinear relationship. To check for nonlinear relationships graphically, we will create a scatterplot and use a simple regression model.

```
[31]: from scipy.stats import pearsonr
def pearsonr_pval(x,y):
    return pearsonr(x,y)[1]

df_all_corr_pValue = data.corr(method=pearsonr_pval)
df_all_corr_pValue = df_all_corr_pValue.unstack().sort_values(kind="quicksort",
    ascending=True).reset_index()
df_all_corr_pValue.rename(columns={"level_0": "Feature 1", "level_1": "Feature_
    2", 0: 'P-Value Correlation'}, inplace=True)
df_all_corr_pValue[df_all_corr_pValue['Feature 1'] == 'annual_salary']
```

```
[31]:
```

	Feature 1	Feature 2	P-Value Correlation
3	annual_salary	base_salary	9.884301e-09
8	annual_salary	balance	4.743560e-02
17	annual_salary	gender	2.789715e-01
21	annual_salary	location_dist	3.323434e-01
34	annual_salary	age	7.184211e-01
38	annual_salary	avg_num_payments	7.646067e-01
48	annual_salary	annual_salary	1.000000e+00

There is inconclusive evidence about the significance of the association between the variables.


2 Plots to see correlations

```
[32]: import seaborn as sns
import matplotlib.pyplot as plt
```

2.1 Base_salary vs Annual Salary

```
[33]: plt.figure(figsize=(8, 6))
sns.regplot("base_salary", "annual_salary", data=data)
```


```
[33]: <matplotlib.axes._subplots.AxesSubplot at 0x1d3f64ed710>
```



2.2 balance vs Annual Salary

```
[34]: plt.figure(figsize=(8, 6))  
sns.regplot("balance", "annual_salary", data=data)
```


```
[34]: <matplotlib.axes._subplots.AxesSubplot at 0x1d3f6567be0>
```



2.3 Age vs Annual Salary

```
[35]: plt.figure(figsize=(8, 6))  
sns.regplot("age", "annual_salary", data=data)
```

```
[35]: <matplotlib.axes._subplots.AxesSubplot at 0x1d3f65b4f98>
```



We have an outlier in “Location distance” feature

```
[36]: data['location_dist'].max()
```

```
[36]: 560.5273886392439
```

3 Model Building - Regression Model

```
[37]: X=data[['age' , 'gender', 'balance', 'location_dist', 'avg_num_payments',
→'base_salary']].values
y=data['annual_salary'].values
```

```
[38]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.2,
→random_state=42)
```

```
[39]: from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(X_train, y_train)
lr.score(X, y)
```

[39]: 0.5231347070503298

3.1 Regression Model - Results

```
[40]: from sklearn.metrics import r2_score
y_predict=lr.predict(X_test)
#R-squared is a statistical measure of how close the data are to the fitted
→regression line.
# It is also known as the coefficient of determination, or the coefficient of
→multiple determination for multiple regression
print('Regression Model, R-squared: ', r2_score(y_test, y_predict))
```

Regression Model, R-squared: 0.4192405895493032

```
[41]: from sklearn.metrics import mean_squared_error
# RMSE - Root Mean Squared Error
print('Regression Model, RMSE', np.sqrt(mean_squared_error(y_test, y_predict)))
```

Regression Model, RMSE 4876.661047214676

3.2 Regression Model Analysis

- The model's R-squared shows that it only explains about 40% of variation in customers' annual salary.
- The RMSE of the model over 20% of the data is near 5000, which indicates somehow the inaccuracy of the model.
- Probably, more data is required to develop a more reliable model.

4 Model Building - Decision Tree

```
[42]: from sklearn.tree import DecisionTreeRegressor
dectree = DecisionTreeRegressor(max_depth=5,random_state=0)
dectree.fit(X_train, y_train)
dectree.score(X_train, y_train)
```

[42]: 0.9630003955678947

4.1 Decision Tree Model - Results

```
[43]: y_predict = dectree.predict(X_test)
print('Decision Tree, R-squared: ', r2_score(y_test, y_predict))
print('Decision Tree, RMSE', np.sqrt(mean_squared_error(y_test, y_predict)))
```

Decision Tree, R-squared: 0.1501796279735662

Decision Tree, RMSE 5899.130613268002

4.2 Decision Tree Analysis

The Decision Tree model got an R-squared value of 0.96 , which is very close to 1, indicating the model can be able to explain the annual salary variability.

However, the model achieved a RMSE of 5899 which is a little bit higher than that of the linear regression model (4876)

5 Final Conclusion

Even though the Decision Tree model is performing better than the Linear Regression model in terms of its R-square value, we can not conclude it will be fine to consider useful as it has a higher RMSE.

This means that the absolute fit of the model is much worse, overall. For that reason, neither of both models should be considered to segment the customers.

Taking in mind that we only have 3 months of data, it would be good to compare the generated models with more data, and also another machine learning technique can be applied in order to see if there is room for improvement.

[]: